

Programming in Go

Matt Holiday
Christmas 2020



File Walk Example

What's the problem?

I want to find duplicate files based on their **content**

What's the problem?

I want to find duplicate files based on their **content**

Use a secure hash, because the names / dates may differ

f088913 2

/Users/mholiday/Dropbox/Emergency/FEMA_P-320_2014_508.pdf

/Users/mholiday/Dropbox/Emergency/nps61-072915-01.pdf

What's the problem?

I want to find duplicate files based on their **content**

Use a secure hash, because the names / dates may differ

```
f088913 2
```

```
  /Users/mholiday/Dropbox/Emergency/FEMA_P-320_2014_508.pdf
```

```
  /Users/mholiday/Dropbox/Emergency/nps61-072915-01.pdf
```

It takes nearly **5 minutes** to comb through my Dropbox folder
(2017 core i7 quad-core MacBook Pro)

Code at <https://github.com/matt4biz/go-class-walk>

Sequential Approach

How it works: Declarations

```
package main

import (
    "crypto/md5"
    "fmt"
    "io"
    "log"
    "os"
    "path/filepath"
)

type pair struct {
    hash, path string
}

type fileList []string
type results  map[string]fileList
```

How it works: Hashing

```
func hashFile(path string) pair {  
    file, err := os.Open(path)  
  
    if err != nil {  
        log.Fatal(err)  
    }  
  
    defer file.Close()  
  
    hash := md5.New() // fast & good enough  
  
    if _, err := io.Copy(hash, file); err != nil {  
        log.Fatal(err)  
    }  
  
    return pair{fmt.Sprintf("%x", hash.Sum(nil)), path}  
}
```


How it works: Searching

```
func searchTree(dir string) (results, error) {  
    hashes := make(results)  
  
    err := filepath.Walk(dir, func(p string, fi os.FileInfo,  
                                   err error) error {  
        // ignore the error parm for now  
  
        if fi.Mode().IsRegular() && fi.Size() > 0 {  
            h := hashFile(p)  
            hashes[h.hash] = append(hashes[h.hash], h.path)  
        }  
  
        return nil  
    })  
  
    return hashes, err  
}
```


How it works: Output

```
func main() {  
    if len(os.Args) < 2 {  
        log.Fatal("Missing parameter, provide dir name!")  
    }  
    if hashes, err := searchTree(os.Args[1]); err == nil {  
        for hash, files := range hashes {  
            if (len(files) > 1) {  
                // we will use just 7 chars like git  
                fmt.Println(hash[len(hash)-7:], len(files))  
  
                for _, file := range files {  
                    fmt.Println("    ", file)  
                }  
            }  
        }  
    }  
}
```

Concurrent Approach #1

A concurrent approach (like map-reduce)

Use a fixed pool of goroutines and a collector and channels



How it works: Collecting the hashes

```
func collectHashes(pairs <-chan pair, result chan<- results) {  
    hashes := make(results)  
  
    for p := range pairs {  
        hashes[p.hash] = append(hashes[p.hash], p.path)  
    }  
  
    result <- hashes  
}
```

How it works: Replacing the processor

```
func processFiles(paths <-chan string, pairs chan<- pair,  
                done chan<- bool) {  
    for path := range paths {  
        pairs <- hashFile(path)  
    }  
  
    done <- true  
}
```

How it works: Replacing the tree walk

```
workers := 2 * runtime.GOMAXPROCS(0)

paths := make(chan string)
pairs := make(chan pair)
done := make(chan bool)
result := make(chan results)

for i := 0; i < workers; i++ {
    go processFiles(paths, pairs, done)
}

// we need another goroutine so we don't block here

go collectHashes(pairs, result)

. . .
```

How it works: Replacing the tree walk

```
err := filepath.Walk(dir, func(p string, fi os.FileInfo,
                                err error) error {
    // again, ignore the error passed in

    if fi.Mode().IsRegular() && fi.Size() > 0 {
        paths <- p
    }

    return nil
})

if err != nil {
    log.Fatal(err)
}

// we must close the paths channel so the workers stop
close(paths)
. . .
```


How it works: Replacing the tree walk

. . .

// wait for all the workers to be done

```
for i := 0; i < workers; i++ {  
    <-done  
}
```

*// by closing pairs we signal that all the hashes
// have been collected; we have to do it here AFTER
// all the workers are done*

```
close(pairs)
```

```
hashes := <-result
```

```
return hashes
```

Evaluation #1

56.11s in the version shown above

Evaluation #1

56.11s in the version shown above

52.76s with a buffer to pairs (buffering pairs keeps the workers working)

Evaluation #1

56.11s in the version shown above

52.76s with a buffer to pairs (buffering pairs keeps the workers working)

51.36s with twice as many workers (32)

Concurrent Approach #2

Another concurrent approach

Add a goroutine for each directory in the tree

Another concurrent approach

Add a goroutine for each directory in the tree

This improves the performance slightly, we're not waiting on paths to be identified

How it works: Parallel tree walk

. . .

```
wg := new(sync.WaitGroup)
```

```
// multi-threaded walk of the directory tree; we need a  
// waitGroup because we don't know how many to wait for
```

```
wg.Add(1)
```

```
err := walkDir(dir, paths, wg)
```

```
if err != nil {  
    log.Fatal(err)  
}
```

```
wg.Wait()  
close(paths)
```

. . .

How it works: Parallel tree walk

```
func walkDir(dir string, paths chan<- string,  
            wg *sync.WaitGroup) error {  
    defer wg.Done()  
  
    visit := func(p string, fi os.FileInfo, err error) error {  
        // ignore the error passed in  
  
        // ignore dir itself to avoid an infinite loop!  
        if fi.Mode().IsDir() && p != dir {  
            wg.Add(1)  
            go walkDir(p, paths)  
            return filepath.SkipDir  
        }  
  
        . . .  
    }
```

How it works: Parallel tree walk

```
. . .  
  
    if fi.Mode().IsRegular() && fi.Size() > 0 {  
        paths <- p  
    }  
  
    return nil  
}  
  
return filepath.Walk(dir, visit)  
}
```

51.14s in the basic version

Evaluation #2

51.14s in the basic version

50.03 adding buffers on all channels to/from workers

Evaluation #2

51.14s in the basic version

50.03 adding buffers on all channels to/from workers

48.75 with twice as many workers

Concurrent Approach #3

The final approach: Goroutines galore!

Use a goroutine for every directory and file hash

The final approach: Goroutines galore!

Use a goroutine for every directory and file hash

What could go wrong?

The final approach: Goroutines galore!

Use a goroutine for every directory and file hash

What could go wrong?

Without some controls, we'll run out of threads!

The final approach: Goroutines galore!

Use a goroutine for every directory and file hash

What could go wrong?

Without some controls, we'll run out of threads!

GOMAXPROCS doesn't limit threads blocked on syscalls (all our disk I/O)

We'll limit the number of **active** goroutines instead (the ones making syscalls)

Channels as counting semaphores

A goroutine can't proceed without *sending* on the channel

Channels as counting semaphores

A goroutine can't proceed without *sending* on the channel

A channel with buffer size N can accept N sends without blocking (with no intervening reads)

Channels as counting semaphores

A goroutine can't proceed without *sending* on the channel

A channel with buffer size N can accept N sends without blocking (with no intervening reads)

The buffer provides a fixed upper bound (unlike a WaitGroup)

Channels as counting semaphores


A goroutine can't proceed without *sending* on the channel

A channel with buffer size N can accept N sends without blocking (with no intervening reads)

The buffer provides a fixed upper bound (unlike a WaitGroup)

One goroutine can start for each one that finishes (each *reads* from the channel when done)

What that looks like



How it works: Limiting goroutines

```
// we don't need a channel for paths or to signal done but  
// we need a buffered channel to act as a counting semaphore
```

```
wg := new(sync.WaitGroup)  
limits := make(chan bool, workers)  
pairs := make(chan pair, workers)  
result := make(chan results)
```

```
go collect(pairs, result)
```

```
. . .
```


How it works: Limiting goroutines

. . .

```
wg.Add(1)  
err := walkDir(dir, pairs, wg, limits)
```

```
if err != nil {  
    log.Fatal(err)  
}
```

```
wg.Wait()  
close(pairs)
```

```
hashes := <-result
```

```
return hashes
```

How it works: Modified processing

```
func processFile(path string, pairs chan<- pair,
                 wg *sync.WaitGroup, limits chan bool) {
    defer wg.Done()

    limits <- true

    defer func() {
        <-limits
    }()

    pairs <- hashFile(path)
}
```

How it works: Modified tree walk

```
func walkDir(dir string, pairs chan<- pair, wg *sync.WaitGroup,
            limits chan bool) error {
    defer wg.Done()

    visit := func(p string, fi os.FileInfo, err error) error {
        // ignore the error passed in

        if fi.Mode().IsDir() && p != dir {
            wg.Add(1)
            go walkDir(p, pairs, wg, limits)
            return filepath.SkipDir
        }

        . . .
    }
```

How it works: Modified tree walk

```
. . .  
  
    if fi.Mode().IsRegular() && fi.Size() > 0 {  
        wg.Add(1)  
        go processFile(p, pairs, wg, limits)  
    }  
  
    return nil  
}  
  
limits <- true  
  
defer func() {  
    <-limits  
}()  
  
return filepath.Walk(dir, visit)  
}
```

Evaluation #3

46.93s using 32 workers was the best time

Evaluation #3

46.93s using 32 workers was the best time

Increasing the `limits` buffer makes the time grow longer due to disk contention

Evaluation #3

46.93s using 32 workers was the best time


Increasing the `limits` buffer makes the time grow longer due to disk contention

Amdahl's law: speedup is limited by the part (not) parallelized

$$S = \frac{1}{1 - p + (p/s)}$$

Here we've managed about $S = 6.25$ on $s = 8$ processors, or about $p = 96\%$ parallel

Evaluation #3



Conclusions

We don't need to limit *goroutines*

We need to limit *contention* for shared resources