

# Programming in Go

---

Matt Holiday  
Christmas 2020



# Benchmarking

---

## Go benchmarks

Go has standard tools and conventions for running benchmarks

Benchmarks live in test files ending with `_test.go`

You run benchmarks with `go test -bench`

Go only runs the `BenchmarkXXX` functions

## Simple example function

```
func Fib(n int, recursive bool) int {  
    switch n {  
    case 0:  
        return 0  
    case 1:  
        return 1  
    default:  
        if recursive {  
            return Fib(n-1, r) + Fib(n-2, r)  
        }  
        a, b := 0, 1  
        for i := 1; i < n; i++ {  
            a, b = b, a+b  
        }  
        return b  
    }  
}
```

## Simple benchmark tests

```
func BenchmarkFib20T(b *testing.B) {  
    for n := 0; n < b.N; n++ {  
        Fib(20, true) // run the Fib function b.N times  
    }  
}
```

```
func BenchmarkFib20F(b *testing.B) {  
    for n := 0; n < b.N; n++ {  
        Fib(20, false) // run the Fib function b.N times  
    }  
}
```

```
$ go test -bench=. ./fib_test.go
```



```
goos: darwin
```

```
goarch: amd64
```

BenchmarkFib20T-16	20851	59634 ns/op
BenchmarkFib20F-16	94855990	12.6 ns/op

## List vs Slice example

A slice of objects beats a list with pointers



## List vs Slice example

```
package main
import "testing"

type node struct {
    v int           // value in the list node
    t *node
}

func insert(i int, h *node) *node {
    t := &node{i, nil}

    if h != nil {
        h.t = t
    }

    return t
}
```

## List vs Slice example

```
func mkList(n int) *node {  
    var h, t *node  
  
    h = insert(0, h)  
    t = insert(1, h)  
    for i := 2; i < n; i ++ {  
        t = insert(i, t)  
    }  
    return h  
}  
  
func sumList(n *node) (i int) {  
    for h := n; h != nil; h = h.t {  
        i += h.v  
    }  
    return  
}
```



## List vs Slice example

```
func mkSlice(n int) []int {  
    r := make([]int, n)  
  
    for i := 0; i < n; i++ {  
        r[i] = i  
    }  
  
    return r  
}  
  
func sumSlice(l []int) (i int) {  
    for _, v := range l {  
        i += v  
    }  
  
    return  
}
```

## List vs Slice example

```
func BenchmarkList(b *testing.B) {  
    for n := 0; n < b.N; n++ {  
        l := mkList(1200)  
        sumList(l)  
    }  
}
```

```
func BenchmarkSlice(b *testing.B) {  
    for n := 0; n < b.N; n++ {  
        l := mkSlice(1200)  
        sumSlice(l)  
    }  
}
```

```
$ go test -bench=. ./list_test.go
```

BenchmarkList-16	35452	33904 ns/op
BenchmarkSlice-16	769028	1555 ns/op

## List vs Slice example

In this example, we'll separate out the cost of making the list

```
func BenchmarkList(b *testing.B) {  
    l := mkList(1200); b.ResetTimer()  
    for n := 0; n < b.N; n++ {  
        sumList(l)  
    }  
}
```

```
func BenchmarkSlice(b *testing.B) {  
    l := mkSlice(1200); b.ResetTimer()  
    for n := 0; n < b.N; n++ {  
        sumSlice(l)  
    }  
}
```

BenchmarkList-16	885607	1243 ns/op
BenchmarkSlice-16	2910057	414 ns/op

## List vs Slice example

In this version, we keep the value separate from the list node

```
type node struct {  
    v *int  
    t *node  
}  
  
func insert(i int, n *node) *node {  
    t := &node{&i, nil}  
  
    if n != nil {  
        n.t = t  
    }  
  
    return t  
}
```

## List vs Slice example

In this version, we keep the value separate from the list node

```
func sumList(n *node) (i int) {  
    for h := n; h != nil; h = h.t {  
        i += *h.v  
    }  
    return  
}
```

BenchmarkList-16	707344	1596 ns/op
BenchmarkSlice-16	2883547	417 ns/op

What if we go back and include the cost of building the list?

BenchmarkList-16	24858	48690 ns/op
BenchmarkSlice-16	662131	1518 ns/op

## List vs Slice example

We can extend the benchmark to include memory allocations using the `-benchmem` flag

This version is with the list building included (since that's where the allocation takes place)

We don't just allocate more, we do it in smaller chunks

```
$ go test -bench=. -benchmem ./list_test.go
```

```
goos: darwin
```


```
goarch: amd64
```

BenchmarkList-16	23977	49766	ns/op	28800	B/op	2400	allocs/op
------------------	-------	-------	-------	-------	------	------	-----------

BenchmarkSlice-16	741272	1515	ns/op	9728	B/op	1	allocs/op
-------------------	--------	------	-------	------	------	---	-----------

## Forwarding example

Calling lots short methods via dynamic dispatch is very expensive



The cost of calling a function should be proportional to the work it does  
(short inline functions vs longer methods with late binding)

## Forwarding example

```
package forward
import ( "math/rand"; "testing" )

const defaultChars = "01234567 . . . mnopqrstuvwxyz"

func randString(length int, charset string) string {
    b := make([]byte, length)
    for i := range b {
        b[i] = charset[rand.Intn(len(charset))]
    }
    return string(b)
}

type forwarder interface {
    forward(string) int
}
```



## Forwarding example

```
type thing1 struct {  
    t forwarder  
}
```

```
func (t1 *thing1) forward(s string) int {  
    return t1.t.forward(s)  
}
```

```
type thing2 struct {  
    t forwarder  
}
```

```
func (t2 *thing2) forward(s string) int {  
    return t2.t.forward(s)  
}
```

```
type thing3 struct {}
```

## Forwarding example

```
func (t3 *thing3) forward(s string) int {  
    return len(s)  
}  
  
func length(s string) int {  
    return len(s)  
}  
  
func BenchmarkDirect(b *testing.B) {  
    r := randString(rand.Intn(24), defaultChars)  
    h := make([]int, b.N)  
  
    b.ResetTimer()  
    for i := 0; i < b.N; i++ {  
        h[i] = length(r)  
    }  
}
```

## Forwarding example

```
func BenchmarkForward(b *testing.B) {  
    r := randString(rand.Intn(24), defaultChars)  
    h := make([]int, b.N)  
  
    var t3 forwarder = &thing3{}  
    var t2 forwarder = &thing2{t3}  
    var t1 forwarder = &thing1{t2}  
  
    b.ResetTimer()  
    for i := 0; i < b.N; i++ {  
        h[i] = t1.forward(r)  
    }  
}
```

```
$ go test -bench=. ./forward_test.go
```

BenchmarkDirect-8	1000000000	0.627 ns/op
BenchmarkForward-8	189176289	6.35 ns/op

## Forwarding example

Let's make one small change

```
//go:noinline  
func length(s string) int {  
    return len(s)  
}
```

```
$ go test -bench=. ./forward_test.go
```

BenchmarkDirect-8	624469180	1.87 ns/op
BenchmarkForward-8	189169784	6.35 ns/op

Presumably that makes both function calls happen out of line

## Forwarding example

Let's make one other change, not so small


```
func BenchmarkForward(b *testing.B) {  
    r := randString(rand.Intn(24), defaultChars)  
    h := make([]int, b.N)  
    b.ResetTimer()  
    for i := 0; i < b.N; i++ {  
        var t3 forwarder = &thing3{}  
        var t2 forwarder = &thing2{t3}  
        var t1 forwarder = &thing1{t2}  
        h[i] = t1.forward(r)  
    }  
}
```

```
$ go test -bench=. ./forward_test.go
```

BenchmarkDirect-8	623482398	1.88 ns/op
BenchmarkForward-8	23797018	49.2 ns/op

## False sharing example

**False sharing:** cores fight over a cache line for *different* variables



## False sharing example

We're going to let CPUs clobber each other's cache

```
import (
    "sync"
    "testing"
)

const (
    nworker = 8
    buffer  = 1024
)

var wg sync.WaitGroup
```

## False sharing example

```
func run() (total int) {  
    cnt := make([]uint64, nworker) // one cache line  
    in := make([]chan int, nworker)  
    for i := 0; i < nworker; i++ {  
        in[i] = make(chan int, buffer)  
        go fill(10000, in[i])  
    }  
    for i := 0; i < nworker; i++ {  
        wg.Add(1)  
        go count(&cnt[i], in[i])  
    }  
    wg.Wait()  
    for _, v := range cnt {  
        total += int(v)  
    }  
    return  
}
```



## False sharing example

```
func count(cnt *uint64, in <-chan int) {  
    // false sharing  
  
    for i := range in {  
        *cnt += uint64(i)  
    }  
  
    wg.Done()  
}  
  
func fill(n int, in chan<- int) {  
    for i := 0; i < n; i++ {  
        in <- i  
    }  
  
    close(in)  
}
```

## False sharing example

```
func BenchmarkShare(b *testing.B) {  
    for i := 0; i < b.N; i++ {  
        run()  
    }  
}
```

We're going to run with different numbers of cores

```
$ go test -bench=. -benchtime=10s -cpu=2,4,8 ./share_test.go
```

goos: darwin

goarch: amd64

BenchmarkShare-2	3630	3280637	ns/op
BenchmarkShare-4	4291	2832584	ns/op
BenchmarkShare-8	4910	2629852	ns/op

And watch it blow up due to false sharing

## False sharing example

Now we're going to write back a local total instead

```
func count(cnt *uint64, in <-chan int) {  
    var total int  
  
    for i := range in {  
        total += i  
    }  
    *cnt = uint64(total)  
    wg.Done()  
}
```

We actually see real improvement with more cores

BenchmarkShare-2	5526	2206546	ns/op
BenchmarkShare-4	9950	1241392	ns/op
BenchmarkShare-8	13222	998438	ns/op

## A few things to consider

Here are some concerns about CPU/memory benchmarking:

- is the data / code available in cache?
- did you hit a garbage collection?
- did virtual memory have to page in/out?
- did branch prediction work the same way?
- did the compiler remove code via optimization?  
(are there side effects in the code?)
- are you running in parallel? how many cores?
- are those cores physical or virtual?
- are you sharing a core with anything else?
- what other processes are sharing the machine?