

# Programming in Go

---

Matt Holiday  
Christmas 2020



# Profiling

---

## Web proxy example

We're going to write a proxy for the “todo” JSON demo server

```
package main

import (
    "encoding/json"
    "fmt"
    "io/ioutil"
    "log"
    "net/http"

    _ "net/http/pprof"
)

const url = "https://jsonplaceholder.typicode.com/todos/"
```

## Web proxy example

```
type todo struct {
    UserID    int    `json:"userID"`
    ID        int    `json:"id"`
    Title     string `json:"title"`
    Completed bool   `json:"completed"`
}

var mark = map[bool]string{
    false: " ",
    true:  "x",
}

func handler(w http.ResponseWriter, r *http.Request) {
    var item todo

    req, _ := http.NewRequest("GET", url+r.URL.Path[1:], nil)
    . . .
}
```

## Web proxy example

. . .

```
tr := &http.Transport{}  
cli := &http.Client{Transport: tr}  
resp, err := cli.Do(req)
```

```
if err != nil {  
    http.Error(w, err.Error(), http.StatusBadGateway)  
}
```

```
if resp.StatusCode != http.StatusOK {  
    http.NotFound(w, r)  
    return  
}
```

```
body, _ := ioutil.ReadAll(resp.Body)  
. . .
```

## Web proxy example

```
. . .  
  
if err := json.Unmarshal(body, &item); err != nil {  
    http.Error(w, err.Error(),  
        http.StatusInternalServerError)  
    return  
}  
  
fmt.Fprintf(w, "[%s] %d - %s\n", mark[item.Completed],  
    item.ID, item.Title)  
}  
  
func main() {  
    http.HandleFunc("/", handler)  
  
    log.Fatal(http.ListenAndServe(":8080", nil))  
}
```

## Running it

We can now open `http://localhost:8080/debug/pprof` which is the default route

If we look at goroutines, we'll see 4 goroutines initially, but as we exercise the server, that number will grow

```
$ curl http://localhost:8080/1  
[ ] 1 - delectus aut autem
```

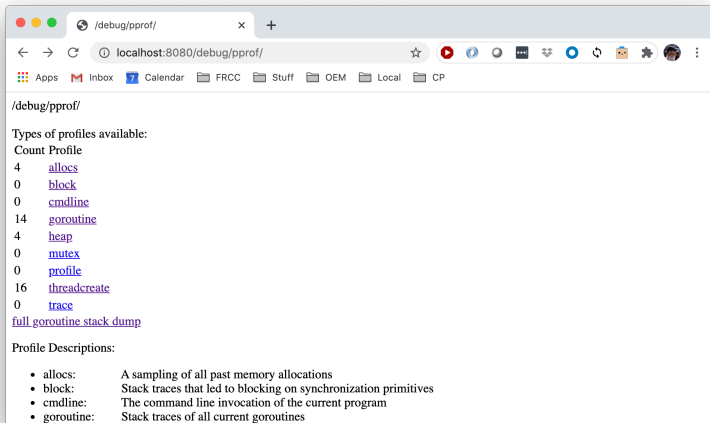
```
$ curl http://localhost:8080/8  
[x] 8 - quo adipisci enim quam ut ab
```

```
$ curl http://localhost:8080/14  
[x] 14 - repellendus sunt dolores architecto voluptatum
```

```
$ curl http://localhost:8080/1133  
404 page not found
```

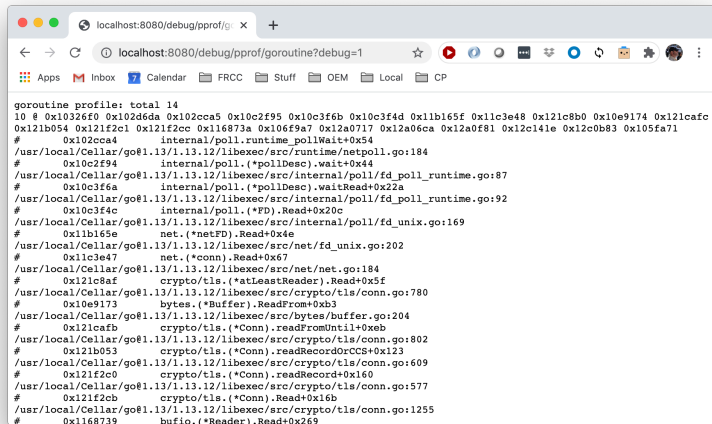
# Main pprof page

This is what including pprof provides for us





In this view we see a list of (hanging) goroutines

A screenshot of a web browser window displaying a goroutine profile. The browser's address bar shows the URL 'localhost:8080/debug/pprof/goroutine?debug=1'. The page content shows a 'goroutine profile: total 14' followed by a list of 14 goroutines. Each entry includes a memory address, a goroutine ID, and the function it is executing. The functions listed are various system calls and library functions like 'runtime\_pollWait', 'pollDesc.wait', 'fd\_poll\_runtime.go', 'Read', 'netFD.Read', 'net.Conn.Read', 'crypto/tls.Conn.readFromUntil', 'crypto/tls.Conn.readRecordOrCCS', and 'bufio.Reader.Read'. The browser's interface includes a standard macOS-style title bar, a navigation bar with back, forward, and refresh buttons, and a toolbar with various icons for bookmarks, extensions, and settings. The browser's address bar also shows the current page title 'localhost:8080/debug/pprof/goroutine?debug=1'.

```
goroutine profile: total 14
10 @ 0x10326f0 0x102d6da 0x102cca5 0x10c2f95 0x10c3f6b 0x10c3f4d 0x11b165f 0x11c3e48 0x121c8b0 0x10e9174 0x121cafc
0x121b054 0x121f2c1 0x121f2cc 0x116873a 0x106f9a7 0x12a0717 0x12a06ca 0x12a0f81 0x12c141e 0x12c0b83 0x105fa71
#      0x102cca4      internal/poll.runtime_pollWait+0x54
/usr/local/Cellar/go@1.13/1.13.12/libexec/src/runtime/netpoll.go:184
#      0x10c2f94      internal/poll.(*pollDesc).wait+0x44
/usr/local/Cellar/go@1.13/1.13.12/libexec/src/internal/poll/fd_poll_runtime.go:87
#      0x10c3f6a      internal/poll.(*pollDesc).waitRead+0x22a
/usr/local/Cellar/go@1.13/1.13.12/libexec/src/internal/poll/fd_poll_runtime.go:92
#      0x10c3f4c      internal/poll.(*FD).Read+0x20c
/usr/local/Cellar/go@1.13/1.13.12/libexec/src/internal/poll/fd_unix.go:169
#      0x11b165e      net.(*netFD).Read+0x4e
/usr/local/Cellar/go@1.13/1.13.12/libexec/src/net/fd_unix.go:202
#      0x11c3e47      net.(*conn).Read+0x67
/usr/local/Cellar/go@1.13/1.13.12/libexec/src/net/net.go:184
#      0x121c8af      crypto/tls.(*atLeastReader).Read+0x5f
/usr/local/Cellar/go@1.13/1.13.12/libexec/src/crypto/tls/conn.go:780
#      0x10e9173      bytes.(*Buffer).ReadFrom+0xb3
/usr/local/Cellar/go@1.13/1.13.12/libexec/src/bytes/buffer.go:204
#      0x121cafb      crypto/tls.(*Conn).readFromUntil+0xeb
/usr/local/Cellar/go@1.13/1.13.12/libexec/src/crypto/tls/conn.go:802
#      0x121b053      crypto/tls.(*Conn).readRecordOrCCS+0x123
/usr/local/Cellar/go@1.13/1.13.12/libexec/src/crypto/tls/conn.go:609
#      0x121f2c0      crypto/tls.(*Conn).readRecord+0x160
/usr/local/Cellar/go@1.13/1.13.12/libexec/src/crypto/tls/conn.go:577
#      0x121f2cb      crypto/tls.(*Conn).Read+0x16b
/usr/local/Cellar/go@1.13/1.13.12/libexec/src/crypto/tls/conn.go:1255
#      0x1168739      bufio.(*Reader).Read+0x269
```

## Fixing it

We'll see that all the goroutines are handling HTTP sockets

And the allocations happen where we read from the response body but forgot to close it

```
// we need this to recover the socket  
  
defer resp.Body.Close()  
  
if resp.StatusCode != http.StatusOK {  
    http.NotFound(w, r)  
    return  
}  
  
body, _ := ioutil.ReadAll(resp.Body)
```

## Prometheus metrics

We can add a Prometheus client which exposes metrics (they can be scraped by hitting /metrics)

```
import ()
    "github.com/prometheus/client_golang/prometheus"
    "github.com/prometheus/client_golang/prometheus/promhttp"
)

var queries = prometheus.NewCounter(prometheus.CounterOpts{
    Name: "all_queries",
    Help: "How many queries we've received.",
})
```

## Prometheus metrics

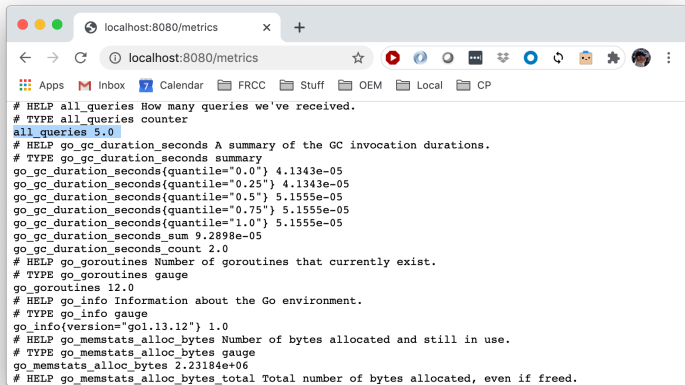
We add a line of code to the end of our handler

```
. . .  
fmt.Fprintf(w, "[%s] %d - %s\n", mark[item.Completed], item.ID, item.Title)  
queries.Inc()  
}
```

And an extra route in our main function

```
func main() {  
    prometheus.MustRegister(queries)  
  
    http.HandleFunc("/", handler)  
    http.Handle("/metrics", promhttp.Handler())  
    . . .  
}
```

We get our metrics as well as a many system metrics



```
# HELP all_queries How many queries we've received.
# TYPE all_queries counter
all_queries 5.0
# HELP go_gc_duration_seconds A summary of the GC invocation durations.
# TYPE go_gc_duration_seconds summary
go_gc_duration_seconds{quantile="0.0"} 4.1343e-05
go_gc_duration_seconds{quantile="0.25"} 4.1343e-05
go_gc_duration_seconds{quantile="0.5"} 5.1555e-05
go_gc_duration_seconds{quantile="0.75"} 5.1555e-05
go_gc_duration_seconds{quantile="1.0"} 5.1555e-05
go_gc_duration_seconds_sum 9.2898e-05
go_gc_duration_seconds_count 2.0
# HELP go_goroutines Number of goroutines that currently exist.
# TYPE go_goroutines gauge
go_goroutines 12.0
# HELP go_info Information about the Go environment.
# TYPE go_info gauge
go_info{version="go1.13.12"} 1.0
# HELP go_memstats_alloc_bytes Number of bytes allocated and still in use.
# TYPE go_memstats_alloc_bytes gauge
go_memstats_alloc_bytes 2.23184e+06
# HELP go_memstats_alloc_bytes_total Total number of bytes allocated, even if freed.
```

## Prometheus metrics

Certain metrics will be a dead giveaway that we're leaking sockets

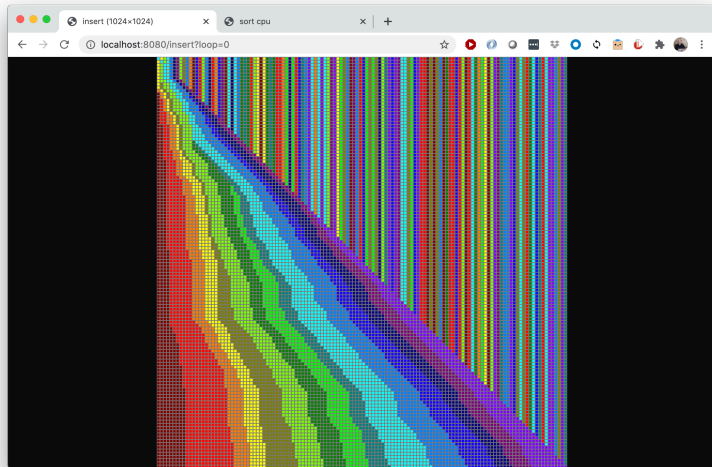
Not only will the goroutine count grow, but also the count of open file descriptors (FDs) *[Linux only — not on macOS]*

```
# HELP go_goroutines Number of goroutines that currently exist.  
# TYPE go_goroutines gauge  
go_goroutines 19.0
```

```
# HELP go_threads Number of OS threads created.  
# TYPE go_threads gauge  
go_threads 17.0
```

```
# HELP process_open_fds Number of open file descriptors.  
# TYPE process_open_fds gauge  
process_open_fds 19.0
```

# Sort animation example



## Sort animation example

The program animates various sort algorithms

- insertion sort: `insert`
- versions of quicksort: `qsort`, `qsortm`, `qsort3`, `qsorti`, `qsortf`

It creates an animated GIF with one frame for each row that changes (each step in the sort until it's done)

It must draw a square of an entry's color outlined in gray

- 1024 rows and columns
- each entry is a square 8 pixels on a side



## How to profile

We must build the program as a binary: `go build .`

We're going to use `pprof` again and hit the endpoint `http://localhost:8080/debug/pprof/profile`

After 30 seconds it will download a profile file

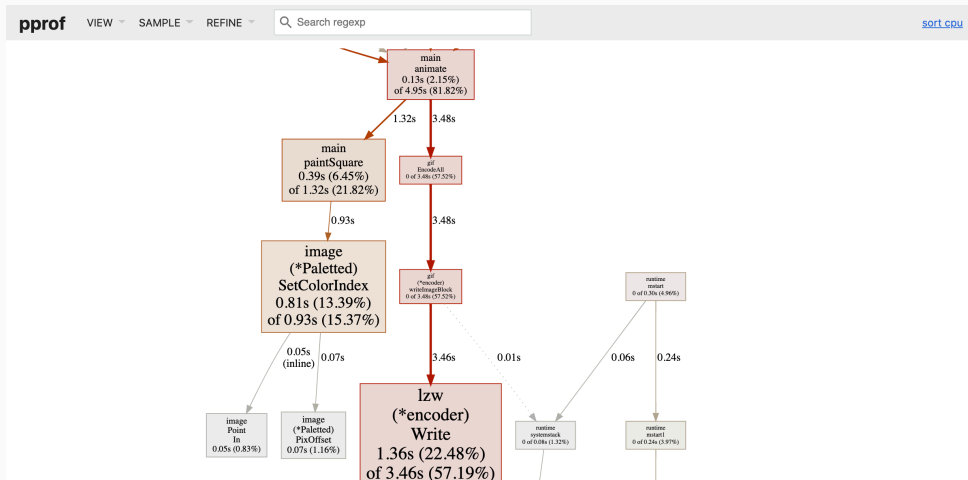
Then we can run the tool: `go tool pprof <binary> <profile-file>`  
in one of three ways

- interactive, with a prompt
- just get the top entries with `-top`
- open a browser with `-http=":6060"`

## CPU flame graph



# CPU profile



# CPU profile

```
$ go tool pprof -top sort profile-slow
```

```
File: sort
```

```
Type: cpu
```

```
Time: Jan 31, 2021 at 8:29am (MST)
```

```
Duration: 30s, Total samples = 6.05s (20.17%)
```

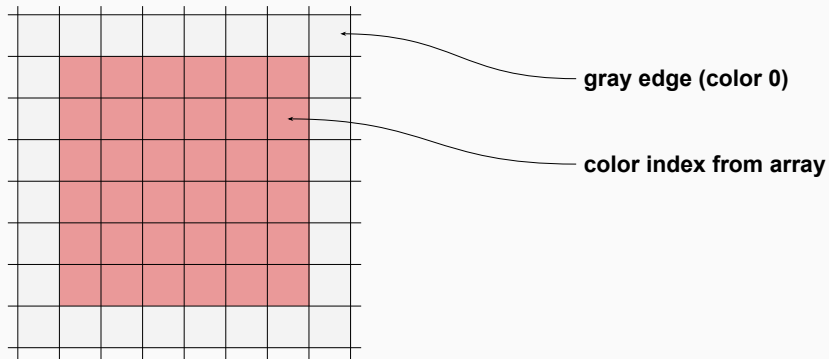
```
Showing nodes accounting for 5.99s, 99.01% of 6.05s total
```

```
Dropped 31 nodes (cum <= 0.03s)
```

flat	flat%	sum%	cum	cum%	
1.68s	27.77%	27.77%	1.68s	27.77%	syscall.syscall
1.36s	22.48%	50.25%	3.46s	57.19%	compress/lzw.(*encoder).Write
0.81s	13.39%	63.64%	0.93s	15.37%	image.(*Paletted).SetColorIndex
0.39s	6.45%	70.08%	1.32s	21.82%	main.paintSquare
0.25s	4.13%	74.21%	0.25s	4.13%	runtime.nanotime1
0.22s	3.64%	77.85%	0.22s	3.64%	runtime.kevent
0.19s	3.14%	80.99%	2.01s	33.22%	compress/lzw.(*encoder).writeLSB
0.19s	3.14%	84.13%	0.41s	6.78%	runtime.netpoll
0.14s	2.31%	86.45%	1.82s	30.08%	image/gif.blockWriter.WriteByte
0.14s	2.31%	88.76%	0.14s	2.31%	runtime.pthread_cond_wait
0.13s	2.15%	90.91%	4.95s	81.82%	main.animate

## Painted squares

We represent each “item” in the array as a square



## Painting each square

We're using a standard library function `SetColorIndex`

```
func paintSquare(i, k int, src []int, img *image.Paletted) {  
    // lay down a square with an outline using the default  
    // color (gray; we deliberately excluded it from the data)  
  
    for x := 0; x < scale; x++ {  
        for y := 0; y < scale; y++ {  
            idx := uint8(src[i])  
  
            if x == 0 || y == 0 || x == scale-1 || y == scale-1 {  
                idx = 0  
            }  
            img.SetColorIndex(i*scale+x, k*scale+y, idx)  
        }  
    }  
}
```

## Setting the color index

In `image/image.go`, we see it does extra work

- checking the point's location
- recalculating the pixel's offset

```
func (p *Paletted) PixOffset(x, y int) int {  
    return (y-p.Rect.Min.Y)*p.Stride + (x-p.Rect.Min.X)*1  
}
```

```
func (p *Paletted) SetColorIndex(x, y int, index uint8) {  
    if !(Point{x, y}.In(p.Rect)) {  
        return  
    }  
    i := p.PixOffset(x, y)  
    p.Pix[i] = index  
}
```

## Optimize the function

There are four things we can do:

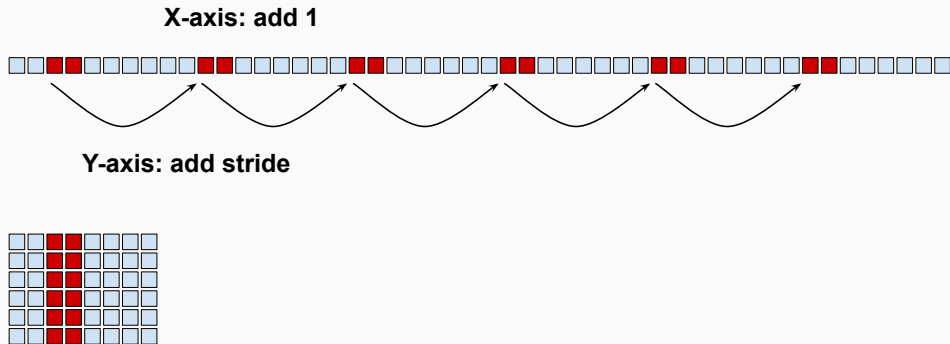
- eliminate the unnecessary check
- move some multiplication out of the loop
- *strength reduction*: replace multiplication with addition *in the loop*
- provided that we reorder the loops (y then x)

We take advantage of the layout of pixels in the image as well as the knowledge that we're filling in a square



# Strength reduction

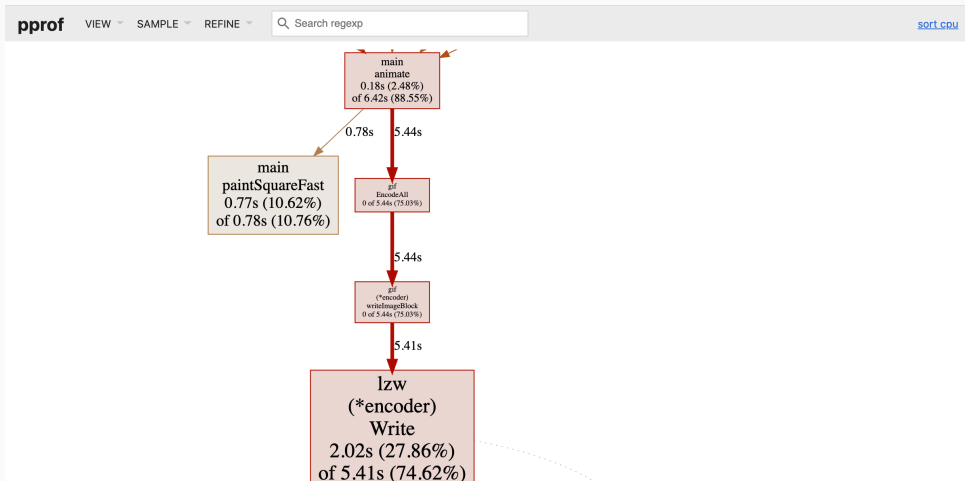
Pixels are kept in a slice, one row after another



## Painting each square #2

```
func paintSquareFast(i, k int, src []int, img *image.Paletted) {  
    ci := uint8(src[i])  
    is, ks := i * scale, k * scale  
    px := (ks-img.Rect.Min.Y)*img.Stride + (is-img.Rect.Min.X)*1  
  
    for y := 0; y < scale; y++ {  
        for x := 0; x < scale; x++ {  
            idx := ci  
  
            if x == 0 || y == 0 || x == scale-1 || y == scale-1 {  
                idx = 0  
            }  
            img.Pix[px+x] = idx  
        }  
        px += img.Stride  
    }  
}
```

# CPU profile #2



## CPU profile #2

```
$ go tool pprof -top sort profile-fast
```

```
File: sort
```

```
Type: cpu
```

```
Time: Jan 31, 2021 at 8:31am (MST)
```

```
Duration: 30s, Total samples = 7.25s (24.17%)
```

```
Showing nodes accounting for 7.11s, 98.07% of 7.25s total
```

```
Dropped 40 nodes (cum <= 0.04s)
```

flat	flat%	sum%	cum	cum%	
2.93s	40.41%	40.41%	2.93s	40.41%	syscall.syscall
2.02s	27.86%	68.28%	5.41s	74.62%	compress/lzw.(*encoder).Write
0.77s	10.62%	78.90%	0.78s	10.76%	main.paintSquareFast
0.35s	4.83%	83.72%	0.35s	4.83%	runtime.nanotime1
0.30s	4.14%	87.86%	3.37s	46.48%	compress/lzw.(*encoder).writeLSB
0.18s	2.48%	90.34%	6.42s	88.55%	main.animate

# CPU profile #2



## Optimize the function, part deux

There are two more things we can do:

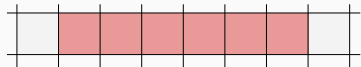
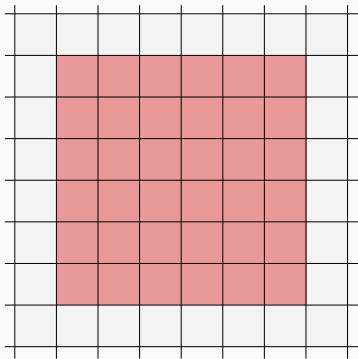
- eliminate the if-then logic in the loop
- reduce slice indexing (& bounds checks) from  $O(n^2)$  to  $O(n)$

We can do this by

- splitting the square into top, middle, bottom sections
- copying a slice of pixels of the correct color into the image

## Painted squares

We only have two kinds of rows to paint



## Painting each square #3

```
func paintSquareFastest(i, k int, src []int, img *image.Paletted) {  
    ci := uint8(src[i])  
    is, ks := i*scale, k*scale  
    px := (ks-img.Rect.Min.Y)*img.Stride + (is-img.Rect.Min.X)*1  
    py := px + (scale-1)*img.Stride  
  
    // create a row for the top and bottom  
  
    rw := make([]uint8, scale)  
  
    // paint the top & bottom rows of the square  
  
    copy(img.Pix[px:px+scale], rw)  
    copy(img.Pix[py:py+scale], rw)  
  
    . . .  
}
```



## Painting each square #3

. . .

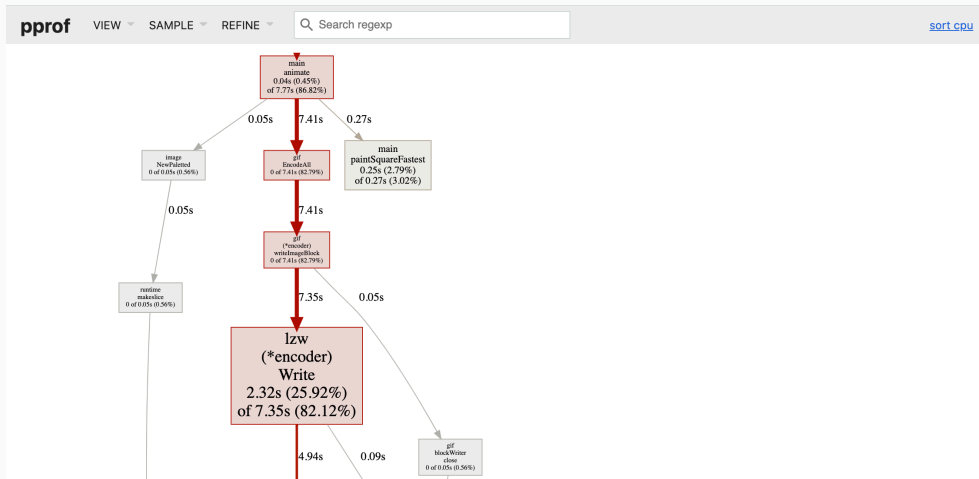
*// fill in the middle part of the row*

```
for x := 1; x < scale-1; x++ {  
    rw[x] = ci  
}
```

*// paint the middle part of the square*

```
for y := 1; y < scale-1; y++ {  
    px += img.Stride  
    copy(img.Pix[px:px+scale], rw)  
}  
}
```

# CPU profile #3



## CPU profile #3

```
$ go tool pprof -top sort /Users/mholiday/Downloads/profile-fastest-3
File: sort
Type: cpu
Time: Jan 31, 2021 at 12:40pm (MST)
Duration: 30.11s, Total samples = 8.95s (29.72%)
Showing nodes accounting for 8.80s, 98.32% of 8.95s total
Dropped 44 nodes (cum <= 0.04s)
      flat  flat%   sum%        cum   cum%   func
    4.52s  50.50%  50.50%    4.53s  50.61%  syscall.syscall
    2.32s  25.92%  76.42%    7.35s  82.12%  compress/lzw.(*encoder).Write
    0.35s   3.91%  80.34%    0.35s   3.91%  runtime.kevent
    0.32s   3.58%  83.91%    0.32s   3.58%  runtime.nanotime1
    0.25s   2.79%  86.70%    0.27s   3.02%  main.paintSquareFastest
    0.24s   2.68%  89.39%    4.94s  55.20%  compress/lzw.(*encoder).writeLSB
    0.15s   1.68%  91.06%    4.70s  52.51%  image/gif.blockWriter.WriteByte

    . . .

    0.04s   0.45%  97.32%    7.77s  86.82%  main.animate
```

# CPU profile #3

